```
!pip install --upgrade google-generativeai
!pip install sentence-transformers faiss-cpu streamlit pyngrok
```

```
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lil
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch>=1.11.0->sentence-t
  Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/py
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/pyth
Requirement already satisfied: smmap<6,>=3.0.1 in /usr/local/lib/python3.1
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.1
Requirement already satisfied: attrs>=22.2.0 in /usr/local/lib/python3.11/
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in /us
Requirement already satisfied: referencing>=0.28.4 in /usr/local/lib/pytho
Requirement already satisfied: rpds-py>=0.7.1 in /usr/local/lib/python3.11/
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-
Downloading faiss_cpu-1.11.0.post1-cp311-cp311-manylinux_2_27_x86_64.manyl
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 31.3/31.3 MB 57.3 MB/s eta 0:0
Downloading streamlit-1.46.1-py3-none-any.whl (10.1 MB)
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 10.1/10.1 MB 106.3 MB/s eta 0:
Downloading pyngrok-7.2.12-py3-none-any.whl (26 kB)
Downloading pydeck-0.9.1-py2.py3-none-any.whl (6.9 MB)
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 6.9/6.9 MB 101.7 MB/s eta 0:00
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 363.4/363.4 MB 1.3 MB/s eta 0:
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 13.8/13.8 MB 72.3 MB/s eta 0:0
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 24.6/24.6 MB 55.4 MB/s eta 0:0
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_6
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 883.7/883.7 kB 45.4 MB/s eta 0
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl (
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 664.8/664.8 MB 2.9 MB/s eta 0:
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl (
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 211.5/211.5 MB 5.6 MB/s eta 0:
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.wh
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 56.3/56.3 MB 11.6 MB/s eta 0:0
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.wh
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 127.9/127.9 MB 7.5 MB/s eta 0:
Downloading nvidia_cusparse_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 207.5/207.5 MB 5.9 MB/s eta 0:
Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.wh
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 21.1/21.1 MB 79.1 MB/s eta 0:00
Installing collected packages: watchdog, pyngrok, nvidia-nvjitlink-cu12, nv
  Attempting uninstall: nvidia-nvjitlink-cu12
```

```
    Found existing installation: nvidia-nvjitlink-cu12 12.5.82
    Uninstalling nvidia-nvjitlink-cu12-12.5.82:
      Successfully uninstalled nvidia-nvjitlink-cu12-12.5.82
  Attempting uninstall: nvidia-curand-cu12
    Found existing installation: nvidia-curand-cu12 10.3.6.82
    Uninstalling nvidia-curand-cu12-10.3.6.82:
      Successfully uninstalled nvidia-curand-cu12-10.3.6.82
  Attempting uninstall: nvidia-cufft-cu12
    Found existing installation: nvidia-cufft-cu12 11.2.3.61
    Uninstalling nvidia-cufft-cu12-11.2.3.61:
```

```
import os
os.environ["GOOGLE_API_KEY"] = "AIzaSyB1fG6M9oFP1ndTg-76AN-iW4U_t2JWuW8"
```

```
code = '''
import os
import google.generativeai as genai
from sentence_transformers import SentenceTransformer
import faiss
import streamlit as st

# Configure Gemini
genai.configure(api_key=os.getenv("GOOGLE_API_KEY"))
model = genai.GenerativeModel("gemini-pro")

# Sample documents
docs = [
    "Python is a versatile programming language.",
    "FAISS enables efficient similarity search over vectors.",
    "Streamlit is used for building web apps in Python quickly.",
    "Gemini is a large language model from Google for smart text generation."
]

# Embedding and indexing
embed_model = SentenceTransformer("all-MiniLM-L6-v2")
doc_embeddings = embed_model.encode(docs)

dimension = doc_embeddings.shape[1]
index = faiss.IndexFlatL2(dimension)
index.add(doc_embeddings)

# Retriever
def retrieve_top_k(query, k=2):
    query_vec = embed_model.encode([query])
    distances, indices = index.search(query_vec, k)
    return [docs[i] for i in indices[0]]
```

```
# Generator
def generate_answer(query, context):
    prompt = f"Context:\\n{context}\\n\\nQuestion: {query}\\nAnswer:"
    response = model.generate_content(prompt)
    return response.text

# UI
st.title("RAG Q&A Chatbot (Gemini)")
query = st.text_input("Ask your question:")

if query:
    context = "\\n".join(retrieve_top_k(query))
    answer = generate_answer(query, context)

    st.subheader("Answer")
    st.write(answer)

    st.subheader("Context")
    st.write(context)
'''

with open("rag_gemini_chatbot.py", "w") as f:
    f.write(code)

print("Chatbot code file created: rag_gemini_chatbot.py")
```

```
✅ Chatbot code file created: rag_gemini_chatbot.py
```

```
!ngrok config add-authtoken "2zvNQ6yvtZxkJivtLC5e3vbxbEd_3J5wMRFaxQJL6KWZZPJYz"
```

```
Authtoken saved to configuration file: /root/.config/ngrok/ngrok.yml
```

```
from pyngrok import ngrok


public_url = ngrok.connect("http://localhost:8501")
print("🔗 Your chatbot is live at:", public_url)

!streamlit run rag_gemini_chatbot.py --server.enableCORS false --server.enableXsr
```

```
•••
    /usr/local/lib/python3.11/dist-packages/google/ai/generativelanguage_v1b
    /generative_service/client.py:835 in generate_content

       832 |   |   self._validate_universe_domain()
```

```
  833 |   |
  834 |   |    # Send the request.
❱ 835 |   |    response = rpc(
  836 |   |    |    request,
  837 |   |    |    retry=retry,
  838 |   |    |    timeout=timeout,
```

/usr/local/lib/python3.11/dist-packages/google/api_core/gapic_v1/**method.**
__call__

```
  128 |   |    if self._compression is not None:
  129 |   |    |    kwargs["compression"] = compression
  130 |   |
❱ 131 |   |    return wrapped_func(*args, **kwargs)
  132
  133
  134 def wrap_method(
```

/usr/local/lib/python3.11/dist-packages/google/api_core/retry/**retry_unary**
retry_wrapped_func

```
  291 |   |   |    sleep_generator = exponential_sleep_generator(
  292 |   |   |    |    self._initial, self._maximum, multiplier=self._mul
  293 |   |   |    )
❱ 294 |   |   |    return retry_target(
  295 |   |   |    |    target,
  296 |   |   |    |    self._predicate,
  297 |   |   |    |    sleep_generator,
```

/usr/local/lib/python3.11/dist-packages/google/api_core/retry/**retry_unary**
retry_target

```
  153 |   |   |    # This function explicitly must deal with broad exceptions
  154 |   |   |    except Exception as exc:
  155 |   |   |    |    # defer to shared logic for handling errors
❱ 156 |   |   |    |    next_sleep = _retry_error_helper(
  157 |   |   |    |    |    exc,
  158 |   |   |    |    |    deadline,
  159 |   |   |    |    |    sleep_iter,
```

/usr/local/lib/python3.11/dist-packages/google/api_core/retry/**retry_base**
_retry_error_helper

```
  211 |   |   |    |    RetryFailureReason.NON_RETRYABLE_ERROR,
  212 |   |   |    |    original_timeout,
  213 |   |   |    )
❱ 214 |   |   |    raise final_exc from source_exc
  215 |   |   if on_error_fn is not None:
  216 |   |   |    on_error_fn(exc)
  217 |   |   # next_sleep is fetched after the on_error callback, to allow
```

/usr/local/lib/python3.11/dist-packages/google/api_core/retry/**retry_unary**
retry_target

Start coding or generate with AI.